

# OptiTree: Advanced Ensemble Strategies for Predictive Accuracy

Preethi Sree Allam

2023-04-17

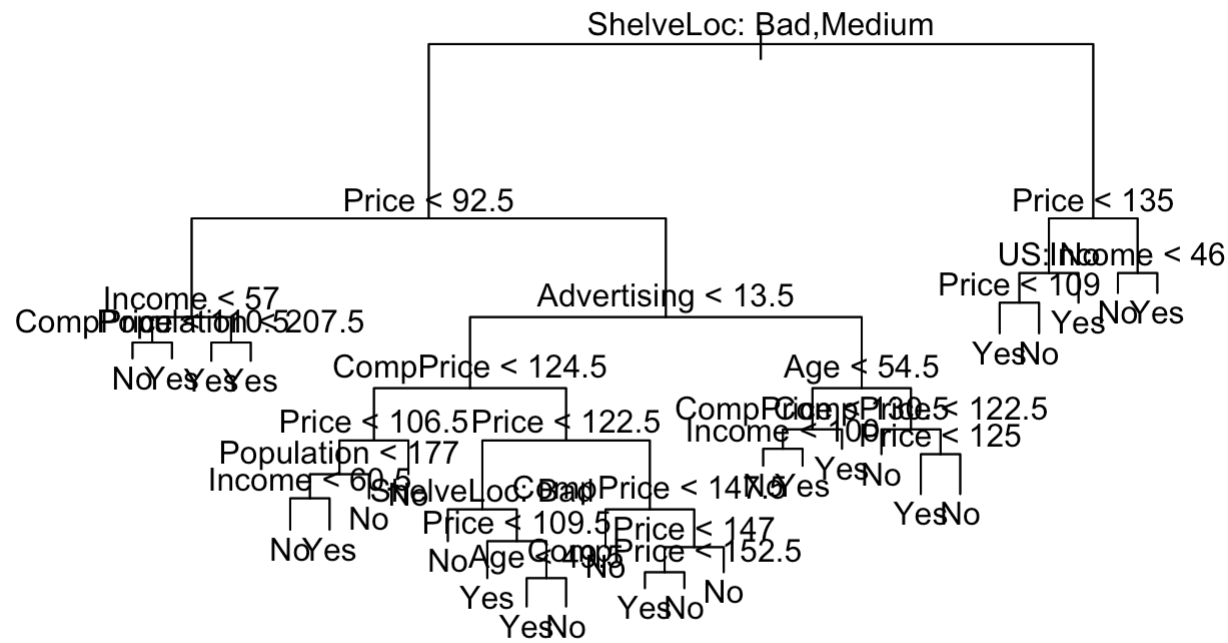
## Decision Trees

### Fitting Classification Trees

```
library(tree)
library(ISLR2)
attach(Carseats)
High <- factor(ifelse(Sales <= 8, "No", "Yes"))
Carseats <- data.frame(Carseats, High)
tree.carseats <- tree(High ~ . - Sales, Carseats)
summary(tree.carseats)
```

```
##
## Classification tree:
## tree(formula = High ~ . - Sales, data = Carseats)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Income" "CompPrice" "Population"
## [6] "Advertising" "Age" "US"
## Number of terminal nodes: 27
## Residual mean deviance: 0.4575 = 170.7 / 373
## Misclassification error rate: 0.09 = 36 / 400
```

```
plot(tree.carseats)
text(tree.carseats, pretty = 0)
```



tree.carseats

```

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 400 541.500 No ( 0.59000 0.41000 )
##    2) ShelfLoc: Bad,Medium 315 390.600 No ( 0.68889 0.31111 )
##      4) Price < 92.5 46 56.530 Yes ( 0.30435 0.69565 )
##        8) Income < 57 10 12.220 No ( 0.70000 0.30000 )
##          16) CompPrice < 110.5 5 0.000 No ( 1.00000 0.00000 ) *
##            17) CompPrice > 110.5 5 6.730 Yes ( 0.40000 0.60000 ) *
##          9) Income > 57 36 35.470 Yes ( 0.19444 0.80556 )
##            18) Population < 207.5 16 21.170 Yes ( 0.37500 0.62500 ) *
##              19) Population > 207.5 20 7.941 Yes ( 0.05000 0.95000 ) *
##        5) Price > 92.5 269 299.800 No ( 0.75465 0.24535 )
##          10) Advertising < 13.5 224 213.200 No ( 0.81696 0.18304 )
##            20) CompPrice < 124.5 96 44.890 No ( 0.93750 0.06250 )
##              40) Price < 106.5 38 33.150 No ( 0.84211 0.15789 )
##                80) Population < 177 12 16.300 No ( 0.58333 0.41667 )
##                  160) Income < 60.5 6 0.000 No ( 1.00000 0.00000 ) *
##                    161) Income > 60.5 6 5.407 Yes ( 0.16667 0.83333 ) *
##                  81) Population > 177 26 8.477 No ( 0.96154 0.03846 ) *
##                41) Price > 106.5 58 0.000 No ( 1.00000 0.00000 ) *
##              21) CompPrice > 124.5 128 150.200 No ( 0.72656 0.27344 )
##                42) Price < 122.5 51 70.680 Yes ( 0.49020 0.50980 )
##                  84) ShelfLoc: Bad 11 6.702 No ( 0.90909 0.09091 ) *
##                    85) ShelfLoc: Medium 40 52.930 Yes ( 0.37500 0.62500 )
##                      170) Price < 109.5 16 7.481 Yes ( 0.06250 0.93750 ) *
##                        171) Price > 109.5 24 32.600 No ( 0.58333 0.41667 )
##                          342) Age < 49.5 13 16.050 Yes ( 0.30769 0.69231 ) *
##                            343) Age > 49.5 11 6.702 No ( 0.90909 0.09091 ) *
##                        43) Price > 122.5 77 55.540 No ( 0.88312 0.11688 )
##                          86) CompPrice < 147.5 58 17.400 No ( 0.96552 0.03448 ) *
##                            87) CompPrice > 147.5 19 25.010 No ( 0.63158 0.36842 )
##                              174) Price < 147 12 16.300 Yes ( 0.41667 0.58333 )
##                                348) CompPrice < 152.5 7 5.742 Yes ( 0.14286 0.85714 ) *
##                                  349) CompPrice > 152.5 5 5.004 No ( 0.80000 0.20000 ) *
##                                    175) Price > 147 7 0.000 No ( 1.00000 0.00000 ) *
##                                11) Advertising > 13.5 45 61.830 Yes ( 0.44444 0.55556 )
##                                  22) Age < 54.5 25 25.020 Yes ( 0.20000 0.80000 )

```

```
##          44) CompPrice < 130.5 14  18.250 Yes ( 0.35714 0.64286 )
##          88) Income < 100 9   12.370 No ( 0.55556 0.44444 ) *
##          89) Income > 100 5    0.000 Yes ( 0.00000 1.00000 ) *
##          45) CompPrice > 130.5 11   0.000 Yes ( 0.00000 1.00000 ) *
##          23) Age > 54.5 20   22.490 No ( 0.75000 0.25000 )
##          46) CompPrice < 122.5 10   0.000 No ( 1.00000 0.00000 ) *
##          47) CompPrice > 122.5 10  13.860 No ( 0.50000 0.50000 )
##          94) Price < 125 5    0.000 Yes ( 0.00000 1.00000 ) *
##          95) Price > 125 5    0.000 No ( 1.00000 0.00000 ) *
##          3) ShelveLoc: Good 85  90.330 Yes ( 0.22353 0.77647 )
##          6) Price < 135 68  49.260 Yes ( 0.11765 0.88235 )
##          12) US: No 17  22.070 Yes ( 0.35294 0.64706 )
##          24) Price < 109 8    0.000 Yes ( 0.00000 1.00000 ) *
##          25) Price > 109 9   11.460 No ( 0.66667 0.33333 ) *
##          13) US: Yes 51  16.880 Yes ( 0.03922 0.96078 ) *
##          7) Price > 135 17  22.070 No ( 0.64706 0.35294 )
##          14) Income < 46 6    0.000 No ( 1.00000 0.00000 ) *
##          15) Income > 46 11  15.160 Yes ( 0.45455 0.54545 ) *
```

```
set.seed(2)
train <- sample(1:nrow(Carseats), 200)
Carseats.test <- Carseats[~train, ]
High.test <- High[~train]
tree.carseats <- tree(High ~ . - Sales, Carseats,
                      subset = train)
tree.pred <- predict(tree.carseats, Carseats.test,
                     type = "class")
table(tree.pred, High.test)
```

```
##          High.test
## tree.pred  No Yes
##          No 104 33
##          Yes 13 50
```



```
set.seed(7)
cv.carseats <- cv.tree(tree.carseats, FUN = prune.misclass)
names(cv.carseats)
```

```
## [1] "size" "dev" "k" "method"
```

```
cv.carseats
```

```
## $size
## [1] 21 19 14 9 8 5 3 2 1
##
## $dev
## [1] 75 75 75 74 82 83 83 85 82
##
## $k
## [1] -Inf 0.0 1.0 1.4 2.0 3.0 4.0 9.0 18.0
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

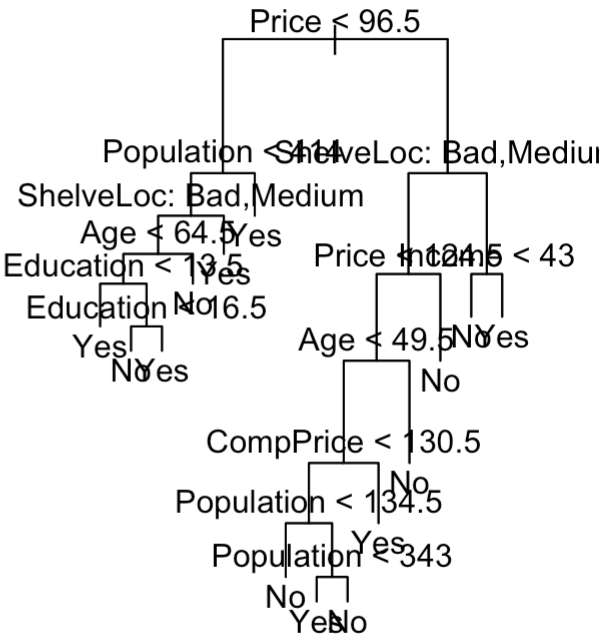
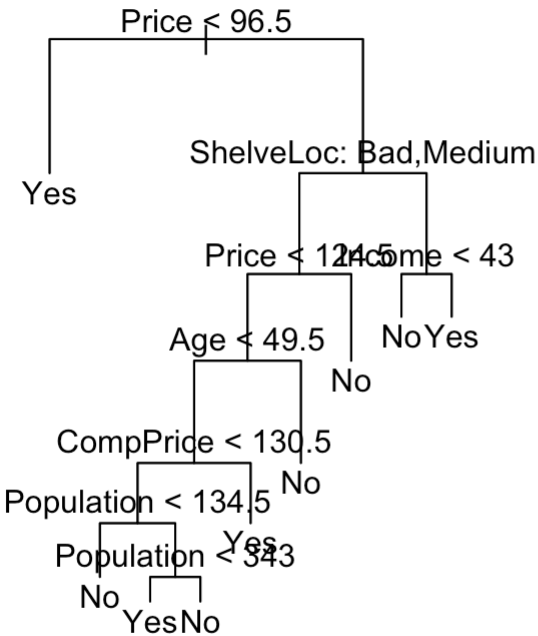
```
par(mfrow = c(1, 2))
plot(cv.carseats$size, cv.carseats$dev, type = "b")
plot(cv.carseats$k, cv.carseats$dev, type = "b")
```



```
prune.carseats <- prune.misclass(tree.carseats, best = 9)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
tree.pred <- predict(prune.carseats, Carseats.test,
                     type = "class")
table(tree.pred, High.test)
```

```
##           High.test
## tree.pred No Yes
##           No  97  25
##           Yes  20  58
```

```
prune.carseats <- prune.misclass(tree.carseats, best = 14)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
```



```
tree.pred <- predict(prune.carseats, Carseats.test,
                    type = "class")
table(tree.pred, High.test)
```

```
##           High.test
## tree.pred  No  Yes
##           No 102  31
##           Yes  15  52
```


## Fitting Regression Trees

```
set.seed(1)
train <- sample(1:nrow(Boston), nrow(Boston) / 2)
tree.boston <- tree(medv ~ ., Boston, subset = train)
summary(tree.boston)
```


```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm"    "lstat" "crim"  "age"
## Number of terminal nodes:  7
## Residual mean deviance:  10.38 = 2555 / 246
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -10.1800 -1.7770  -0.1775   0.0000   1.9230  16.5800
```

```
plot(tree.boston)
text(tree.boston, pretty = 0)
```







```
cv.boston <- cv.tree(tree.boston)
plot(cv.boston$size, cv.boston$dev, type = "b")
```



```
prune.boston <- prune.tree(tree.boston, best = 5)
plot(prune.boston)
text(prune.boston, pretty = 0)
```



```
yhat <- predict(tree.boston, newdata = Boston[-train, ])  
boston.test <- Boston[-train, "medv"]  
plot(yhat, boston.test)  
abline(0, 1)
```



```
mean((yhat - boston.test)^2)
```

```
## [1] 35.28688
```

## Bagging and Random Forests

```
library(randomForest)
```


```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1)
bag.boston <- randomForest(medv ~ ., data = Boston,
subset = train, mtry = 12, importance = TRUE)
bag.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = Boston, mtry = 12, importance = TRUE,      subset = train)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 12
##
##              Mean of squared residuals: 11.40162
##              % Var explained: 85.17
```

```
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])
plot(yhat.bag, boston.test)
abline(0, 1)
```



```
mean((yhat.bag - boston.test)^2)
```

```
## [1] 23.41916
```

```
bag.boston <- randomForest(medv ~ ., data = Boston,  
subset = train, mtry = 12, ntree = 25)  
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])  
mean((yhat.bag - boston.test)^2)
```

```
## [1] 25.75055
```

```
set.seed(1)
rf.boston <- randomForest(medv ~ ., data = Boston,
  subset = train, mtry = 6, importance = TRUE)
yhat.rf <- predict(rf.boston, newdata = Boston[-train, ])
mean((yhat.rf - boston.test)^2)
```

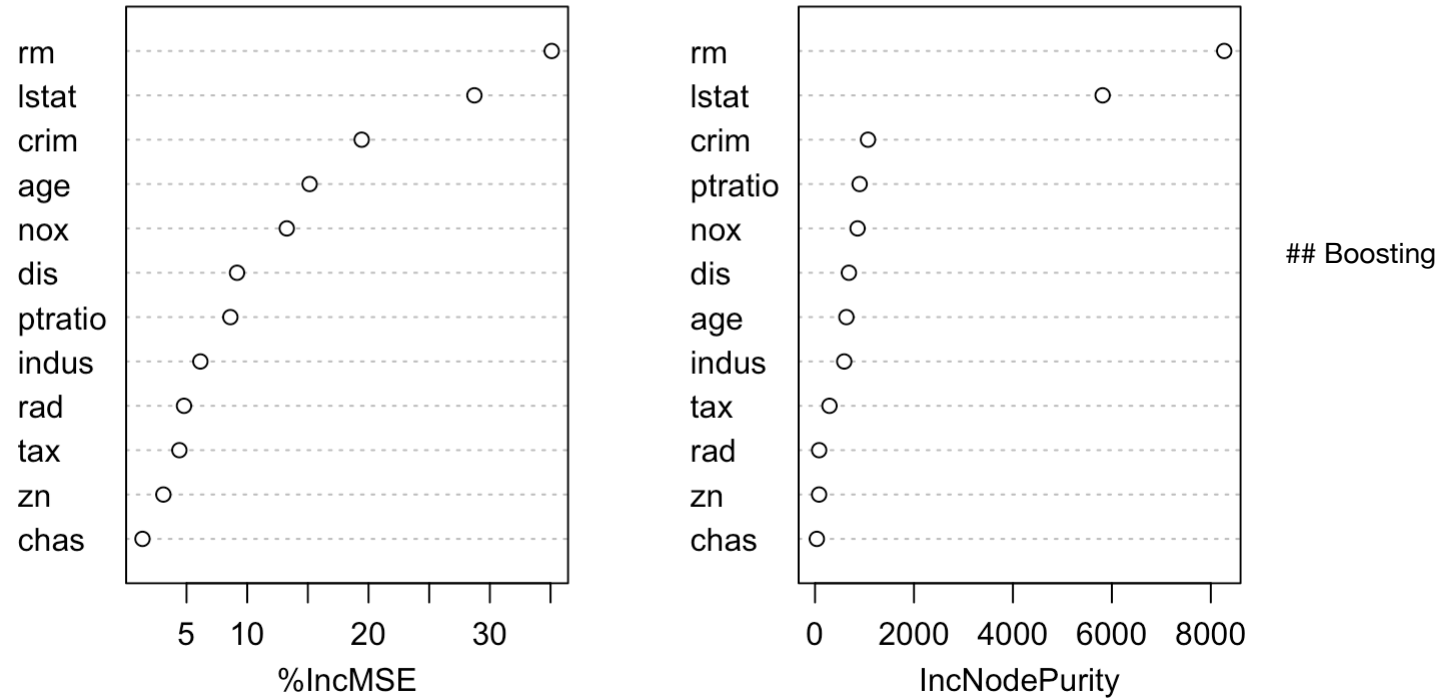
```
## [1] 20.06644
```

```
importance(rf.boston)
```

```
##           %IncMSE IncNodePurity
## crim      19.435587    1070.42307
## zn         3.091630      82.19257
## indus      6.140529    590.09536
## chas       1.370310     36.70356
## nox       13.263466    859.97091
## rm        35.094741   8270.33906
## age       15.144821    634.31220
## dis        9.163776    684.87953
## rad        4.793720     83.18719
## tax        4.410714    292.20949
## ptratio    8.612780    902.20190
## lstat     28.725343   5813.04833
```

```
varImpPlot(rf.boston)
```

rf.boston



```
library(gbm)


## Warning: package 'gbm' was built under R version 4.3.3

## Loaded gbm 2.2.2

## This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com/gbm-developers/gbm3
```




```
set.seed(1)
boost.boston <- gbm(medv ~ ., data = Boston[train, ],
distribution = "gaussian", n.trees = 5000,
interaction.depth = 4)
summary(boost.boston)
```




```
##          var      rel.inf
## rm          rm 44.48249588
## lstat      lstat 32.70281223
## crim       crim  4.85109954
## dis        dis  4.48693083
## nox        nox  3.75222394
## age        age  3.19769210
## ptratio    ptratio 2.81354826
## tax        tax  1.54417603
## indus      indus  1.03384666
## rad        rad  0.87625748
## zn         zn   0.16220479
## chas       chas  0.09671228
```

```
plot(boost.boston, i = "rm")
```



```
plot(boost.boston, i = "lstat")
```



```
yhat.boost <- predict(boost.boston,  
newdata = Boston[-train, ], n.trees = 5000)  
mean((yhat.boost - boston.test)^2)
```

```
## [1] 18.39057
```

```
boost.boston <- gbm(medv ~ ., data = Boston[train, ],  
distribution = "gaussian", n.trees = 5000,  
interaction.depth = 4, shrinkage = 0.2, verbose = F)  
yhat.boost <- predict(boost.boston,  
newdata = Boston[-train, ], n.trees = 5000)  
mean((yhat.boost - boston.test)^2)
```

```
## [1] 16.54778
```

## Bayesian Additive Regression Trees

```
library(BART)
```

```
## Warning: package 'BART' was built under R version 4.3.3
```

```
## Loading required package: nlme
```

```
## Loading required package: survival
```

```
x <- Boston[, 1:12]  
y <- Boston[, "medv"]  
xtrain <- x[train, ]  
ytrain <- y[train]  
xtest <- x[-train, ]  
ytest <- y[-train]  
set.seed(1)  
bartfit <- gbart(xtrain, ytrain, x.test = xtest)
```

```
## *****Calling gbart: type=1
## *****Data:
## data:n,p,np: 253, 12, 253
## y1,yn: 0.213439, -5.486561
## x1,x[n*p]: 0.109590, 20.080000
## xp1,xp[np*p]: 0.027310, 7.880000
## *****Number of Trees: 200
## *****Number of Cut Points: 100 ... 100
## *****burn,nd,thin: 100,1000,1
## *****Prior:beta,alpha,tau,nu,lambda,offset: 2,0.95,0.795495,3,3.71636,21.7866
## *****sigma: 4.367914
## *****w (weights): 1.000000 ... 1.000000
## *****Dirichlet:sparse,theta,omega,a,b,rho,augment: 0,0,1,0.5,1,12,0
## *****printevery: 100
##
## MCMC
## done 0 (out of 1100)
## done 100 (out of 1100)
## done 200 (out of 1100)
## done 300 (out of 1100)
## done 400 (out of 1100)
## done 500 (out of 1100)
## done 600 (out of 1100)
## done 700 (out of 1100)
## done 800 (out of 1100)
## done 900 (out of 1100)
## done 1000 (out of 1100)
## time: 2s
## trcnt,tecnt: 1000,1000
```

```
yhat.bart <- bartfit$yhat.test.mean
mean((ytest - yhat.bart)^2)
```

```
## [1] 15.91912
```

```
ord <- order(bartfit$varcount.mean, decreasing = T)
bartfit$varcount.mean[ord]
```

```
##      nox   lstat    rad      rm    tax ptratio   chas    age   indus    zn
## 22.973 21.653 21.638 20.725 20.021 19.615 19.283 19.278 19.073 15.576
##      dis    crim
## 13.800 11.607
```


```
library(tree)
library(randomForest)
library(MASS)
```

```
##
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:ISLR2':
##
## Boston
```

```
set.seed(1)
subset<-sample(1:nrow(Boston),nrow(Boston)*0.7)
Boston.train<-Boston[subset,-14]
Boston.test<-Boston[-subset,-14]
y.train<-Boston[subset,14]
y.test<-Boston[-subset,14]
rfmodel1<-randomForest(Boston.train,y=y.train,xtest = Boston.test,ytest = y.test,ntree=500,mtry=ncol(Boston.train))
rfmodel2<-randomForest(Boston.train,y.train,xtest = Boston.test,ytest = y.test,ntree=500,mtry=(ncol(Boston.train))/2)
rfmodel3<-randomForest(Boston.train,y.train,xtest = Boston.test,ytest = y.test,ntree=500,mtry=(ncol(Boston.train))^(0.5))
rfmodel4<-randomForest(Boston.train,y.train,xtest = Boston.test,ytest = y.test,ntree=500,mtry=(ncol(Boston.train))/3)
rfmodel5<-randomForest(Boston.train,y.train,xtest = Boston.test,ytest = y.test,ntree=500,mtry=(ncol(Boston.train))/4)
plot(1:500,rfmodel1$test$mse,col="red",type="l",xlab = "Number of Trees",ylab = "Test MSE",ylim = c(10,25))
lines(1:500,rfmodel2$test$mse, col="orange",type="l")
lines(1:500,rfmodel3$test$mse, col="green",type="l")
lines(1:500,rfmodel4$test$mse, col="blue",type="l")
lines(1:500,rfmodel5$test$mse, col="black",type="l")
legend("topright",c("m=p/13","m=p/2","m=sqrt(p)","m=p/3","m=p/4"),col=c("red","orange","green","blue","black"),ce
x=0.5,lty=1)
```





We see that Test MSE decreases with the increase in number of trees. It stabilizes after certain number of trees and no further improvement is seen.

```
library(ISLR)
```

```
##  
## Attaching package: 'ISLR'
```

```
## The following object is masked _by_ '.GlobalEnv':  
##  
##   Carseats
```

```
## The following objects are masked from 'package:ISLR2':  
##  
##   Auto, Credit
```

```
attach(Carseats)
```

```
## The following object is masked _by_ .GlobalEnv:  
##  
##   High
```

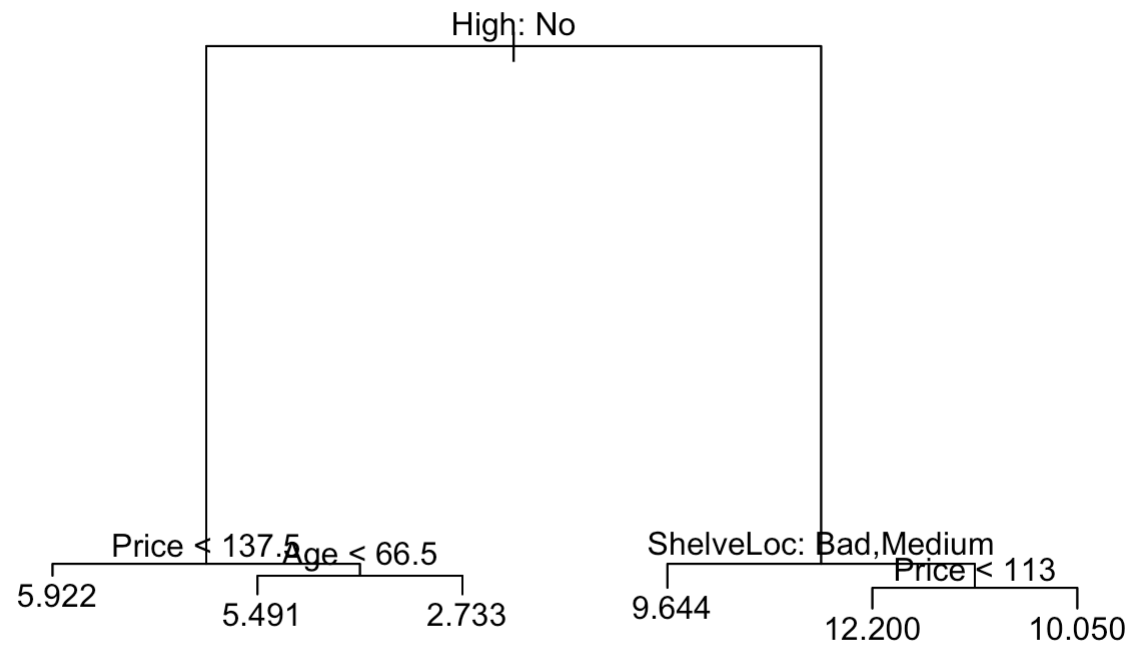
```
## The following objects are masked from Carseats (pos = 10):  
##  
##   Advertising, Age, CompPrice, Education, Income, Population, Price,  
##   Sales, ShelfLoc, Urban, US
```

```
set.seed(1)  
subset<-sample(nrow(Carseats),nrow(Carseats)*0.7)  
car.train<-Carseats[subset,]  
car.test<-Carseats[-subset,]
```

```
library(tree)  
car.model.train<-tree(Sales~.,car.train)  
summary(car.model.train)
```

```
##  
## Regression tree:  
## tree(formula = Sales ~ ., data = car.train)  
## Variables actually used in tree construction:  
## [1] "High"      "Price"     "Age"       "ShelveLoc"  
## Number of terminal nodes: 6  
## Residual mean deviance:  2.081 = 570.2 / 274  
## Distribution of residuals:  
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## -4.502  -0.964  -0.037   0.000   1.048   3.716
```

```
plot(car.model.train)  
text(car.model.train,pretty=0)
```

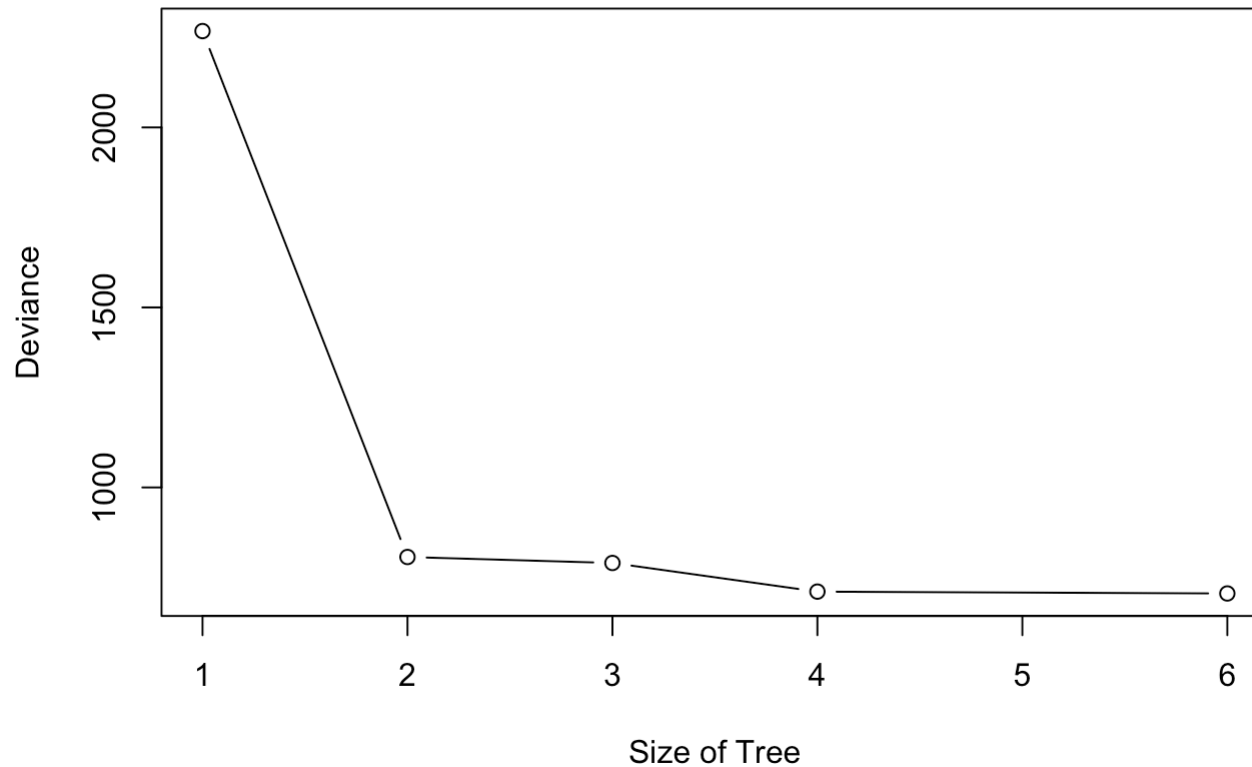


```
tree.prediction<-predict(car.model.train,newdata=car.test)
tree.mse<-mean((car.test$Sales-tree.prediction)^2)
tree.mse
```

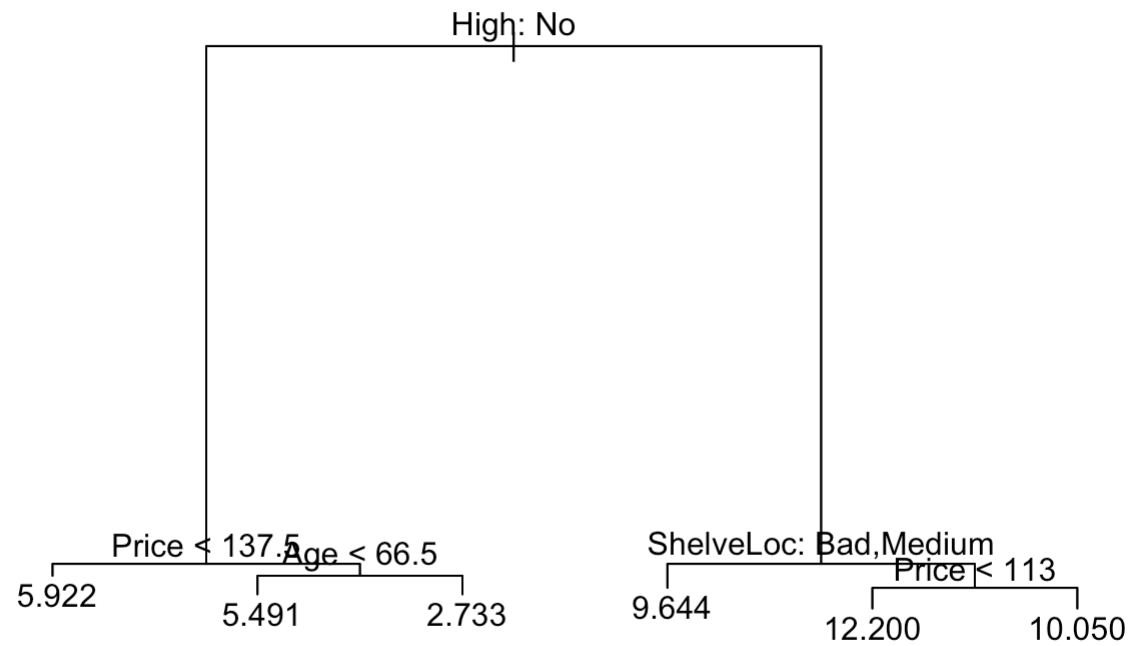
```
## [1] 2.98086
```

## The Test MSE for full grown Tree is recorded as 5.288

```
set.seed(1)
cv.car<-cv.tree(car.model.train)
plot(cv.car$size,cv.car$dev,xlab = "Size of Tree",ylab = "Deviance",type = "b")
```



```
prune.car<-prune.tree(car.model.train,best=6)
plot(prune.car)
text(prune.car,pretty=0)
```



```
prune.predict<-predict(prune.car,car.test)
mean((prune.predict-car.test$Sales)^2)
```

```
## [1] 2.98086
```

**For the pruned tree we get MSE as 5.454**

```
bag.car<-randomForest(Sales~.,car.train,importance=TRUE,mtry=13)
```

```
## Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid
## range
```

```
importance(bag.car)
```

```
##           %IncMSE IncNodePurity
## CompPrice  16.9808558    94.501792
## Income    -2.3357235    58.429280
## Advertising 8.0939012    45.420033
## Population -4.0682805    54.279255
## Price      36.3535209   220.436675
## ShelfLoc   33.8350571   125.105257
## Age        11.4439739   100.602908
## Education  -0.3099822    36.484335
## Urban      -1.3278151     8.833596
## US         1.4225721     6.988115
## High      109.4175231  1470.147822
```

```
bag.car.predict<-predict(bag.car,car.test)
mean((bag.car.predict-car.test$Sales)^2)
```

```
## [1] 2.041283
```

**We use randomForest with  $m=p=13$  total number of predictors which is equivalent to bagging, The Test Set MSE obtained is 2.324. It has further reduced compared to single pruned tree. Thus Bagging helped reducing the MSE, We can see that Price & ShelfLoc are the two most important variables chosen by Bagging model**

```
rf.car<-randomForest(Sales~.,car.train,importance=TRUE,mtry=sqrt(13))
importance(rf.car)
```

```
##           %IncMSE IncNodePurity
## CompPrice 10.4245307    109.82698
## Income   -2.1469464     91.34584
## Advertising 5.8209621    77.60837
## Population -2.7845501    72.96286
## Price     23.3821519   310.36233
## ShelveLoc 22.8678599   284.92725
## Age       6.0439606   138.37473
## Education -0.5646153    55.22101
## Urban     -4.1408498    13.12375
## US        3.2582924    15.01993
## High      69.2591486  1020.45181
```

```
rf.car.predict<-predict(rf.car,car.test)
mean((rf.car.predict-car.test$Sales)^2)
```

```
## [1] 2.041314
```


**Using Random Forest the MSE increased compared to bagging, The important variables chosen by Random Forest are the same as the one chosen by Bagging. Random Forest avoids correlated trees and hence is expected to perform better than Bagging. Here the case is not true. Further tuning of Random Forest model should be tried: Full Grown Tree MSE: 5.288; Pruned Tree MSE: 5.454; Bagging Model MSE: 2.324; Random Forest MSE: 2.518**

```
attach(Hitters)
Hitters<-na.omit(Hitters)
Hitters$Salary<-log(Hitters$Salary)
```


```
subset<-1:200
hitters.train<-Hitters[subset,]
hitters.test<-Hitters[-subset,]
```



```
library(gbm)
set.seed(1)
powerss<-seq(-2,0,by=0.1)
lambdas<-10^powerss
train.error<-rep(NA,length(lambdas))
for (i in 1:length(lambdas)){
  hitters.gbm<-gbm(Salary~.,hitters.train,distribution = "gaussian",n.trees=1000,shrinkage=lambdas[i])
  hitters.train.pred<-predict(hitters.gbm,hitters.train,n.trees=1000)
  train.error[i]<-mean((hitters.train.pred-hitters.train$Salary)^2)
}
plot(lambdas,train.error,type="b",xlab="Shrinkage Value(lambda)",ylab="Training MSE")
```



```
set.seed(1)
test.error<-rep(NA,length(lambdas))
for (i in 1:length(lambdas)){
  hitters.gbm<-gbm(Salary~.,hitters.train,distribution = "gaussian",n.trees=1000,shrinkage=lambdas[i])
  hitters.test.pred<-predict(hitters.gbm,hitters.test,n.trees=1000)
  test.error[i]<-mean((hitters.test.pred-hitters.test$Salary)^2)
}
plot(lambdas,test.error,type="b",xlab="Shrinkage Value(lambda)",ylab="Test MSE")
```



```
hitters.gbm.testerror<-min(test.error)
hitters.gbm.testerror
```

```
## [1] 0.255455
```

**The Minimum Test MSE obtained by boosting is 0.26.**

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
lm<-lm(Salary~.,hitters.train)
hitters.predict.lm<-predict(lm,hitters.test)
hitters.lm.test.mse<-mean((hitters.predict.lm-hitters.test$Salary)^2)
hitters.lm.test.mse
```

```
## [1] 0.4917959
```

```
x<-model.matrix(Salary~.,hitters.train)
x.test<-model.matrix(Salary ~ . , hitters.test)
y<-hitters.train$Salary
hitters.ridge<-glmnet(x,y,alpha=0)
hitters.ridge.predict<-predict(hitters.ridge,s=0.01,x.test)
hitters.ridge.test.mse<-mean((hitters.ridge.predict-hitters.test$Salary)^2)
hitters.ridge.test.mse
```

```
## [1] 0.4570283
```

```
x<-model.matrix(Salary~.,hitters.train)
x.test<-model.matrix(Salary ~ . , hitters.test)
y<-hitters.train$Salary
hitters.lasso<-glmnet(x,y,alpha=1)
hitters.lasso.predict<-predict(hitters.lasso,s=0.01,x.test)
hitters.lasso.test.mse<-mean((hitters.lasso.predict-hitters.test$Salary)^2)
hitters.lasso.test.mse
```

```
## [1] 0.4700537
```

**We have Test MSE for different methods as summarized below. It can be seen Boosting gives least Test MSE among the 4 models**


**Least Square Regression Full Model Test MSE: 0.49**

**Ridge Regression Model Test MSE: 0.46**

**Lasso Regression Model Test MSE: 0.47**

```
boost.hitters<-gbm(Salary~.,data=hitters.train,distribution = "gaussian",n.trees = 1000,shrinkage=lambdas[which.m
in(test.error)])

summary(boost.hitters)
```



```
##          var      rel.inf
## CAtBat    CAtBat 30.0391455
## Years     Years  7.1722320
## CWalks    CWalks 6.6962390
## PutOuts   PutOuts 6.1919523
## Walks     Walks  6.0430398
## CRuns     CRuns  5.8184446
## CHmRun    CHmRun 5.7355580
## CRBI      CRBI   5.6678930
## Hits      Hits   5.5180489
## HmRun     HmRun  3.7274075
## Assists   Assists 3.6165621
## AtBat     AtBat  2.8770937
## RBI       RBI    2.8062318
## CHits     CHits  2.8030774
## Errors    Errors 2.3509666
## Runs      Runs   1.5093746
## Division  Division 0.6614964
## NewLeague NewLeague 0.5632541
## League    League  0.2019828
```

**We find that CAtbat is the most important variable**

```
set.seed(1)
hitters.bagging<-randomForest(Salary~.,hitters.train,mtry=19,importance=TRUE)
hitters.bagging.predict<-predict(hitters.bagging,hitters.test)
hitters.bagging.test.mse<-mean((hitters.bagging.predict-hitters.test$Salary)^2)
hitters.bagging.test.mse
```

```
## [1] 0.2301184
```

**The Test set MSE for Bagging is 0.23 .This is lower than the Test set MSE obtained for Boosting which was 0.26**